



# ANZ synthesized transaction (Task 2)

By: Atefeh Gholipour

# Project Description

	A	B	C	D	E	F	G	H	I	J	K	L	M	N	O	P	Q	R	S	T	U	V	W
1	status	card_pres	bpay	account	currency	long_lat	txn_descr	merchant	merchant	first_name	balance	date	gender	age	merchant	merchant	extraction	amount	transaction_id	country	customer_id	merchant	movement
2	authorize	1		ACC-1598	AUD	153.41 -27	POS	81c48296-73be-44a7	Diana		35.39	8/1/2018	F	26	Ashmore	QLD	2018-08-0	16.25	a623070bf	Australia	CUS-2487	153.38 -27	debit
3	authorize	0		ACC-1598	AUD	153.41 -27	SALES-PO	830a451c-316e-4a6a	Diana		21.2	8/1/2018	F	26	Sydney	NSW	2018-08-0	14.19	13270a2a5	Australia	CUS-2487	151.21 -33	debit
4	authorize	1		ACC-1222	AUD	151.23 -33	POS	835c231d-8cdf-4e96	Michael		5.71	8/1/2018	M	38	Sydney	NSW	2018-08-0	6.42	feb79e7e	Australia	CUS-2142	151.21 -33	debit
5	authorize	1		ACC-1037	AUD	153.10 -27	SALES-PO	48514682-c78a-4a88	Rhonda		2117.22	8/1/2018	F	40	Buderim	QLD	2018-08-0	40.9	2698170d	Australia	CUS-1614	153.05 -26	debit
6	authorize	1		ACC-1598	AUD	153.41 -27	SALES-PO	b4e02c10-0852-4273	Diana		17.95	8/1/2018	F	26	Mermaid	QLD	2018-08-0	3.25	329adf798	Australia	CUS-2487	153.44 -28	debit

## Objective:

- Build a simple regression model to predict the annual salary for each customer

## Dataset:

- Contains synthesised transaction of 3 months for 100 customers including inter bank, payment, phone bank, POS, sales-POS and salary transactions.

7): `dataset.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12043 entries, 0 to 12042
Data columns (total 23 columns):
status                12043 non-null object
card_present_flag     7717 non-null float64
bpay_biller_code      885 non-null object
account               12043 non-null object
currency              12043 non-null object
long_lat              12043 non-null object
txn_description       12043 non-null object
merchant_id           7717 non-null object
merchant_code         883 non-null float64
first_name            12043 non-null object
balance               12043 non-null float64
date                  12043 non-null datetime64[ns]
gender                12043 non-null object
age                   12043 non-null int64
merchant_suburb       7717 non-null object
merchant_state        7717 non-null object
extraction            12043 non-null object
amount                12043 non-null float64
transaction_id        12043 non-null object
country               12043 non-null object
customer_id           12043 non-null object
merchant_long_lat     7717 non-null object
movement              12043 non-null object
dtypes: datetime64[ns](1), float64(4), int64(1), object(17)
memory usage: 2.1+ MB
```

# Analysis Steps:

## Step 1: Preliminary analysis and feature selection

- Drop some irrelevant features
- Check for null data (there is no null in remaining dataset)

```
[10]: # drop irrelevant columns

df = dataset.drop(['status', 'card_present_flag', 'bpay_biller_code',
↳ 'account', 'currency', 'merchant_id',
↳ 'merchant_code', 'first_name', 'date', 'merchant_suburb',
↳ 'merchant_state', 'extraction',
↳ 'transaction_id', 'country', 'merchant_long_lat',
↳ 'movement'], axis = 1)
```

```
[11]: df.head()
```

```
[11]:      long_lat txn_description  balance gender  age  amount  customer_id
0  153.41 -27.95          POS    35.39     F   26   16.25  CUS-2487424745
1  153.41 -27.95    SALES-POS    21.20     F   26   14.19  CUS-2487424745
2  151.23 -33.94          POS     5.71     M   38    6.42  CUS-2142601169
3  153.10 -27.66    SALES-POS   2117.22     F   40   40.90  CUS-1614226872
4  153.41 -27.95    SALES-POS    17.95     F   26    3.25  CUS-2487424745
```


```
[12]: # check for nulls
df.isnull().sum()
```

```
[12]: long_lat      0
      txn_description  0
      balance      0
      gender      0
      age         0
      amount      0
      customer_id  0
      dtype: int64
```

# Analysis Steps:

## Step 2: Data Extraction (annual salary)

- Annual salary should be extracted from txn\_description column
- Distribution of the salary of the customers are shown as:



- There is no data for 16-08-2018.
- For finding annual salary from 3 month payment three:

```
[20]: # Calculate annual salary
      # if Salary_Count >= 12 then payment is weekly
      # if Salary_Count <= 5 then payment is monthly
      # if other value then payment is fortnightly

[21]: df_CI_S ['Annual_Salary'] = 0
      for i in range(0, len(df_CI_S)):
          if int(df_CI_S.Salary_Count[i]) >= 12:
              df_CI_S.Annual_Salary[i] = df_CI_S ['PAY/SALARY'][i]/(df_CI_S.
↳Salary_Count[i]) / 7 * 356
          elif int(df_CI_S.Salary_Count[i]) <= 5:
              df_CI_S.Annual_Salary[i] = df_CI_S ['PAY/SALARY'][i]/(df_CI_S.
↳Salary_Count[i]) * 12
          else:
              df_CI_S.Annual_Salary[i] = df_CI_S ['PAY/SALARY'][i]/(df_CI_S.
↳Salary_Count[i]) / 14 * 356

      # all transaction multiply 4 to obtain for one year

df_CI_S [['INTER BANK', 'PAYMENT', 'PHONE BANK', 'POS',
          'SALES-POS']] = 4 * df_CI_S [['INTER BANK', 'PAYMENT', 'PHONE BANK', 'POS',
          'SALES-POS']]
```

# Analysis Steps:

## Step 2: Data Extraction (customers' location)

- State of each customer are found based on latitude and longitude information.
- There is one out-of-range latitude. So, remove it (row number 2036).

```
# find state of each customer base on longitude and latitude information

import geopy
from geopy.geocoders import Nominatim
def LoctoState(LOC):
    locator = Nominatim(user_agent="myGeocoder")
    coordinates = LOC
    location = locator.reverse(coordinates, exactly_one = True, timeout = 10)
    address = location.raw['address']
    state = address.get('state','')
    return state
#print (LoctoState( "-27.51,153.03"))

df_S_B['state'] = ''
for i in range (0, len(df_S_B)):
    df_S_B['state'][i] = LoctoState(df_S_B['lat_long'][i])
```

# Analysis Steps:

## Step 2: Data Extraction

- Attribute 'total spend' is defined which includes inter bank, payment, phone bank, POS and sales-POS.
- Initial balance of each account has been added.
- The final dataset for analysis is as below:


	age	gender	INTER BANK	PAYMENT	PHONE BANK	POS	SALES-POS	Annual_Salary	balance	total spend	state
0	53	F	0	5184	2184	2992.04	4251.40	49355	463.96	14611.44	Queensland
1	21	M	4004	15828	0	2425.48	13477.80	90999	2335.35	35735.28	Western Australia
2	28	M	1080	3408	0	5425.88	12132.28	48734	823.53	22046.16	Victoria
3	34	F	1000	10388	0	8249.24	7293.76	87036	1726.28	26931.00	New South Wales
4	34	F	3068	12068	0	9222.60	10539.84	99266	12529.59	34898.44	South Australia

# Analysis Steps:

## Step 3: Data visualisation

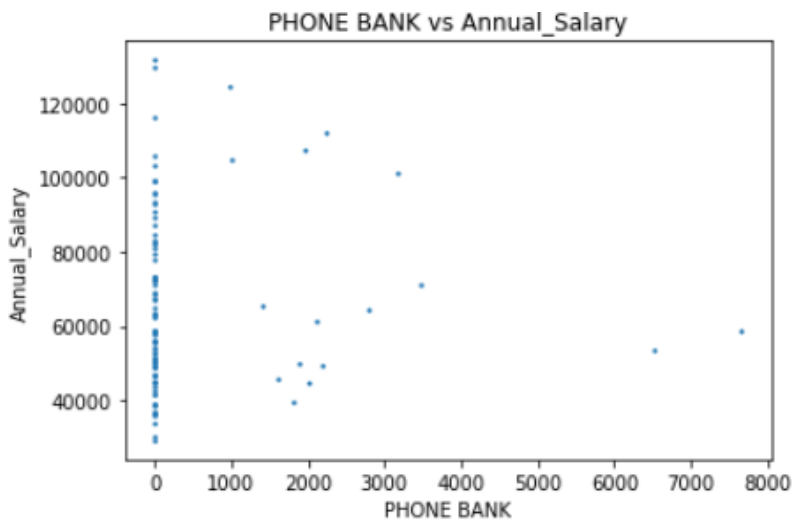


- Distribution of features

```
[25]: # Histogram of Data visualisation  
df_F.hist(rwidth = 0.9)  
plt.tight_layout()
```



# Analysis Steps:





## Step 3: Data visualisation (continuous features)





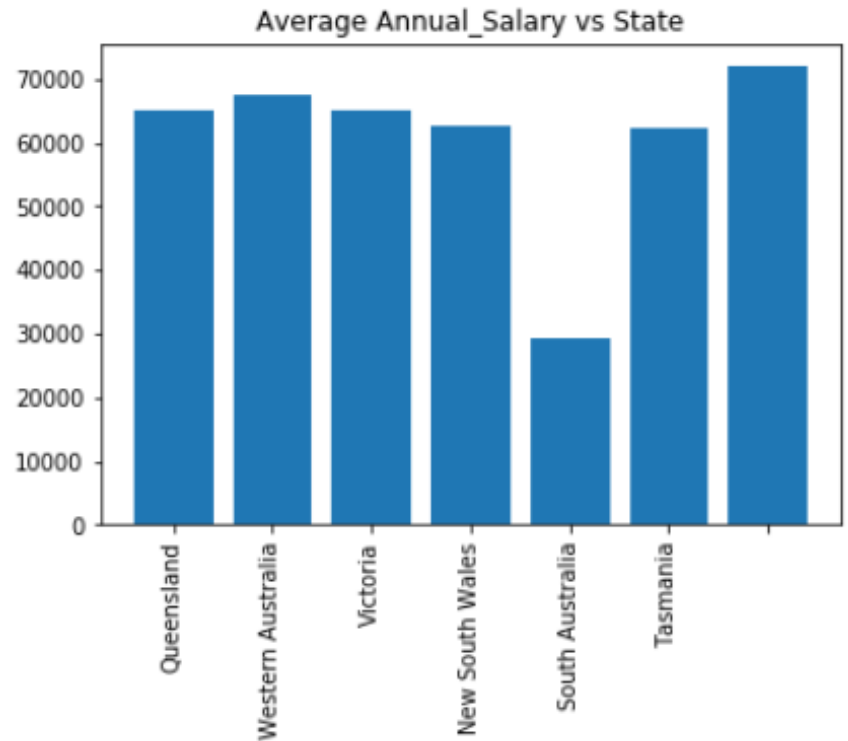
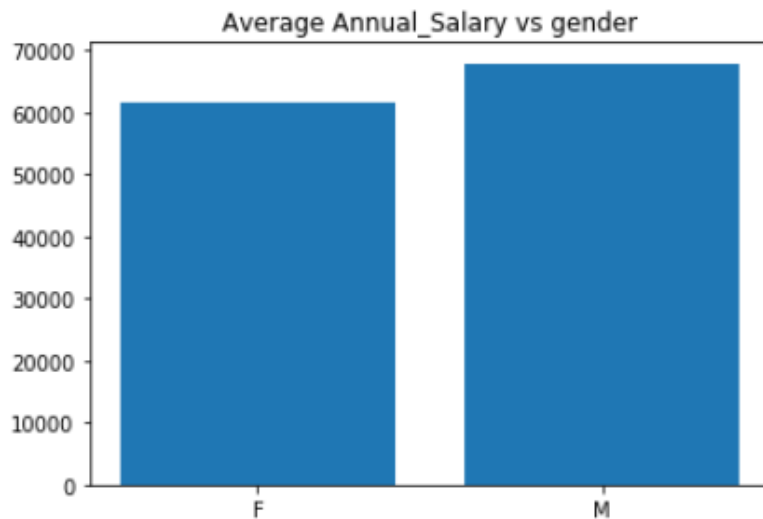
# Analysis Steps:

## Step 3: Data visualisation (continuous features)



# Analysis Steps:

## Step 3: Data visualisation (categorical features)



# Analysis Steps:

## Step 4: Feature selection (correlation calculation)

- Drop features with lower correlation

```
# Check linearity using correlation coefficient matrix
correlation = df_F[['Annual_Salary', 'age', 'INTER BANK', 'PAYMENT', 'PHONE BANK',
                    'POS', 'SALES-POS',
                    'balance', 'total spend']].corr()
print(correlation)
```

	Annual_Salary	age	INTER BANK	PAYMENT	PHONE BANK	\
Annual_Salary	1.000000	-0.061377	0.352362	0.639631	0.033414	
age	-0.061377	1.000000	-0.099233	0.026884	0.103961	
INTER BANK	0.352362	-0.099233	1.000000	0.087386	-0.081680	
PAYMENT	0.639631	0.026884	0.087386	1.000000	-0.132095	
PHONE BANK	0.033414	0.103961	-0.081680	-0.132095	1.000000	
POS	-0.086938	-0.036929	0.181437	-0.123618	-0.052313	
SALES-POS	0.100400	-0.139284	0.158792	0.121610	0.007496	
balance	0.110321	0.237992	0.211241	0.018268	0.026537	
total spend	0.371378	-0.086176	0.476295	0.416710	0.015633	

	POS	SALES-POS	balance	total spend
Annual_Salary	-0.086938	0.100400	0.110321	0.371378
age	-0.036929	-0.139284	0.237992	-0.086176
INTER BANK	0.181437	0.158792	0.211241	0.476295
PAYMENT	-0.123618	0.121610	0.018268	0.416710
PHONE BANK	-0.052313	0.007496	0.026537	0.015633
POS	1.000000	0.418105	-0.000239	0.689232
SALES-POS	0.418105	1.000000	-0.153300	0.756790
balance	-0.000239	-0.153300	1.000000	0.002437
total spend	0.689232	0.756790	0.002437	1.000000

# Analysis Steps:

## Step 4: Feature selection (f\_regression)

- get\_dummies for categorical variables and StandardScaler for continuous variables
- Use f\_regression to calculate F-score and P-value
- Feature with lower F\_score has been removed

```
# feature selection
from sklearn.feature_selection import f_regression
result = f_regression(X,Y)
f_score = result [0]
p_value = result [1]

columns = list(X.columns)
print(" Feature ", " F_score ", " P_value ")
print(" -----", "-----", "-----")
for i in range(0, len(columns)):
    f1= f_score[i]
    p1= p_value[i]
    print(" ", columns[i].ljust(20), " ", f1, " ", p1)
```

Feature	F_score	P_value
-----	-----	-----
INTER BANK	13.077682186681905	0.00047643710935370927
PAYMENT	68.01558513144904	7.985674327460962e-13
SALES-POS	0.9104035219992912	0.34237914303848604
balance	1.2633164571082347	0.26379816386664495
total spend	14.845159693260255	0.0002095941070946692
gender_M	1.7205564326030651	0.19271752801591077
state_New South Wales	0.3735928480105821	0.5424821000548368
state_Queensland	0.0010899400623479772	0.9737310705621645
state_South Australia	0.07316966142052807	0.7873511721337136
state_Tasmania	2.350633544897465	0.12848730409545903
state_Victoria	0.8304095476594371	0.3644132122518978
state_Western Australia	1.337959517193049	0.25023486278067525

# Analysis Steps:

## Step 4: Linear Regression

- Final dataset is as follows:

	gender	INTER BANK	PAYMENT	Annual_Salary	total spend
0	F	0	5184	49355	14611.44
1	M	4004	15828	90999	35735.28
2	M	1080	3408	48734	22046.16
3	F	1000	10388	87036	26931.00
4	F	3068	12068	99266	34898.44

- After splitting data to test and train datasets and fitting the model:

```
score_test = LR.score(X_test, Y_test)
RMSE = math.sqrt(mean_squared_error(Y_test, Y_Predicted))

print('r2 score for Linear Regression is:' ,score_test)
print('RMSE for Linear Regression is:' ,RMSE)
```

```
r2 score for Linear Regression is: 0.6173311697006681
RMSE for Linear Regression is: 0.6155336134400503
```

- However some method such as KFold can slightly improve the r2\_value.

```
# KFold
from sklearn.model_selection import KFold
folds = KFold(n_splits = 5, shuffle = True, random_state = 100)
scores = cross_val_score(LR, X_train, Y_train, scoring='r2', cv=folds)
scores

array([ 0.40019539,  0.68358956,  0.58267536, -0.19175429, -0.09101695])
```

# Summary

- Even with using feature selection method, the selected feature can not leads to a good model and  $r^2$ \_value is not high enough ( $r^2$ \_score = 0.61).
- As the number of customer was very low, obtaining the model with low accuracy was predictable.

# ANZ\_regression

January 29, 2021

```
[97]: # import libraries
import pandas as pd
import matplotlib.pyplot as plt
import numpy as np
import math
```

```
[98]: # read dataset
dataset = pd.read_excel("ANZ synthesised transaction dataset.xlsx")
dataset.head()
```

```
[98]:      status  card_present_flag  bpayer_biller_code      account currency \
0  authorized                1.0                NaN  ACC-1598451071      AUD
1  authorized                0.0                NaN  ACC-1598451071      AUD
2  authorized                1.0                NaN  ACC-1222300524      AUD
3  authorized                1.0                NaN  ACC-1037050564      AUD
4  authorized                1.0                NaN  ACC-1598451071      AUD
```

```
      long_lat  txn_description      merchant_id \
0  153.41 -27.95      POS  81c48296-73be-44a7-befa-d053f48ce7cd
1  153.41 -27.95  SALES-POS  830a451c-316e-4a6a-bf25-e37caedca49e
2  151.23 -33.94      POS  835c231d-8cdf-4e96-859d-e9d571760cf0
3  153.10 -27.66  SALES-POS  48514682-c78a-4a88-b0da-2d6302e64673
4  153.41 -27.95  SALES-POS  b4e02c10-0852-4273-b8fd-7b3395e32eb0
```

```
      merchant_code  first_name  ...  age  merchant_suburb  merchant_state \
0                NaN      Diana  ...   26      Ashmore      QLD
1                NaN      Diana  ...   26      Sydney      NSW
2                NaN  Michael  ...   38      Sydney      NSW
3                NaN  Rhonda  ...   40      Buderim      QLD
4                NaN      Diana  ...   26  Mermaid Beach      QLD
```

```
      extraction  amount      transaction_id \
0  2018-08-01T01:01:15.000+0000  16.25  a623070bfead4541a6b0fff8a09e706c
1  2018-08-01T01:13:45.000+0000  14.19  13270a2a902145da9db4c951e04b51b9
2  2018-08-01T01:26:15.000+0000   6.42  feb79e7ecd7048a5a36ec889d1a94270
3  2018-08-01T01:38:45.000+0000  40.90  2698170da3704fd981b15e64a006079e
4  2018-08-01T01:51:15.000+0000   3.25  329adf79878c4cf0aeb4188b4691c266
```

	country	customer_id	merchant_long_lat	movement
0	Australia	CUS-2487424745	153.38 -27.99	debit
1	Australia	CUS-2487424745	151.21 -33.87	debit
2	Australia	CUS-2142601169	151.21 -33.87	debit
3	Australia	CUS-1614226872	153.05 -26.68	debit
4	Australia	CUS-2487424745	153.44 -28.06	debit

[5 rows x 23 columns]

[99]: `dataset.info()`

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 12043 entries, 0 to 12042
Data columns (total 23 columns):
status                12043 non-null object
card_present_flag     7717 non-null float64
bpay_biller_code      885 non-null object
account               12043 non-null object
currency              12043 non-null object
long_lat              12043 non-null object
txn_description        12043 non-null object
merchant_id           7717 non-null object
merchant_code         883 non-null float64
first_name            12043 non-null object
balance               12043 non-null float64
date                  12043 non-null datetime64[ns]
gender                12043 non-null object
age                   12043 non-null int64
merchant_suburb       7717 non-null object
merchant_state        7717 non-null object
extraction            12043 non-null object
amount                12043 non-null float64
transaction_id        12043 non-null object
country               12043 non-null object
customer_id           12043 non-null object
merchant_long_lat     7717 non-null object
movement              12043 non-null object
dtypes: datetime64[ns](1), float64(4), int64(1), object(17)
memory usage: 2.1+ MB
```

[100]: `#dataset.columns`

[101]: `#print(dataset.date.dtype)`  
*# There is no data for 16-08-2018 so ignore this month*  
`#sub_dataset = dataset[(dataset['date'] >= '2018-09-01')]`  
`#sub_dataset = dataset`



```
[102]: # drop irrelevant columns

df = dataset.drop(['status', 'card_present_flag', 'bpay_biller_code',
→ 'account', 'currency', 'merchant_id',
→ 'merchant_code', 'first_name', 'date', 'merchant_suburb',
→ 'merchant_state', 'extraction',
→ 'transaction_id', 'country', 'merchant_long_lat',
→ 'movement'], axis = 1)

[103]: df.head()
```

	long_lat	txn_description	balance	gender	age	amount	customer_id
0	153.41 -27.95	POS	35.39	F	26	16.25	CUS-2487424745
1	153.41 -27.95	SALES-POS	21.20	F	26	14.19	CUS-2487424745
2	151.23 -33.94	POS	5.71	M	38	6.42	CUS-2142601169
3	153.10 -27.66	SALES-POS	2117.22	F	40	40.90	CUS-1614226872
4	153.41 -27.95	SALES-POS	17.95	F	26	3.25	CUS-2487424745

```
[104]: # check for nulls
df.isnull().sum()

[104]: long_lat      0
txn_description    0
balance            0
gender             0
age               0
amount            0
customer_id       0
dtype: int64

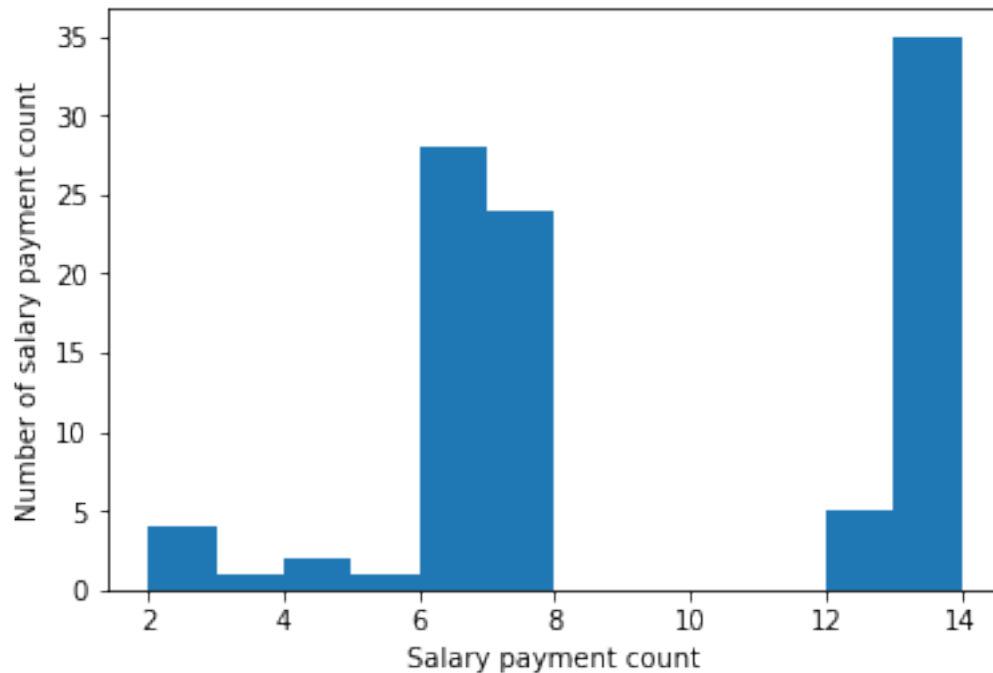
[105]: # Count the number of PAY/SALARY for each customer

df_salary = dataset[['customer_id', 'txn_description', 'amount', 'date']]
df_salary = df_salary[(df['txn_description'] == 'PAY/SALARY')]

#df_salary_sort = df_salary.sort_values(by = ['customer_id', 'date'], ascending
→ = False)
#df_salary_sort

df_salary_count = df_salary.pivot_table(index=['customer_id'], aggfunc='size')
type(df_salary_count)
df_count_S = pd.DataFrame(df_salary_count)
df_count_S = df_count_S.reset_index()
df_count_S.rename(columns = {0 : 'Salary_Count'}, inplace = True)
#type(df_count)
plt.hist(df_salary_count, bins = 12)
plt.ylabel('Number of salary payment count')
plt.xlabel('Salary payment count')
#df_count
```

[105]: Text(0.5, 0, 'Salary payment count')



```
[106]: # pivot table to find unique customers and their salary and other kind of
        ↳ transactions
df_CI = pd.pivot_table(df, values = 'amount',
                        index = ['customer_id', 'age', 'long_lat', 'gender' ],
                        columns = ['txn_description'], aggfunc = np.sum,
                        ↳fill_value = 0)
df_CI.head()
```

```
[106]: txn_description      INTER BANK  PAY/SALARY  PAYMENT \
customer_id  age long_lat  gender
CUS-1005756958  53  153.03 -27.51 F          0    12616.11    1296
CUS-1117979751  21  115.81 -31.82 M       1001    25050.55    3957
CUS-1140341822  28  144.97 -37.42 M        270    11499.06     852
CUS-1147642491  34  151.04 -33.77 F        250    22248.07    2597
CUS-1196156254  34  138.52 -35.01 F        767    27326.11    3017
```

```
txn_description      PHONE BANK      POS  SALES-POS
customer_id  age long_lat  gender
CUS-1005756958  53  153.03 -27.51 F          546    748.01    1062.85
CUS-1117979751  21  115.81 -31.82 M          0    606.37    3369.45
CUS-1140341822  28  144.97 -37.42 M          0   1356.47    3033.07
CUS-1147642491  34  151.04 -33.77 F          0   2062.31    1823.44
CUS-1196156254  34  138.52 -35.01 F          0   2305.65    2634.96
```

```
[107]: #df_CI.columns
```

```
[108]: #df_CI.info()
```

```
[109]: # Turn Pandas Multi-Index into column
```

```
df_CI = df_CI.reset_index()
df_CI.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 100 entries, 0 to 99
Data columns (total 10 columns):
customer_id    100 non-null object
age            100 non-null int64
long_lat       100 non-null object
gender         100 non-null object
INTER BANK     100 non-null int64
PAY/SALARY     100 non-null float64
PAYMENT        100 non-null int64
PHONE BANK     100 non-null int64
POS            100 non-null float64
SALES-POS      100 non-null float64
dtypes: float64(3), int64(4), object(3)
memory usage: 7.9+ KB
```

```
[110]: df_CI.head()
```

```
[110]: txn_description    customer_id  age    long_lat  gender  INTER BANK  \
0      CUS-1005756958    53  153.03 -27.51    F           0
1      CUS-1117979751    21  115.81 -31.82    M        1001
2      CUS-1140341822    28  144.97 -37.42    M        270
3      CUS-1147642491    34  151.04 -33.77    F        250
4      CUS-1196156254    34  138.52 -35.01    F        767
```

```
txn_description  PAY/SALARY  PAYMENT  PHONE BANK    POS  SALES-POS
0      12616.11      1296      546  748.01  1062.85
1      25050.55      3957      0  606.37  3369.45
2      11499.06      852      0  1356.47  3033.07
3      22248.07      2597      0  2062.31  1823.44
4      27326.11      3017      0  2305.65  2634.96
```

```
[111]: # Add Salary Count to dataset
```

```
df_CI_S = df_CI.merge(df_count_S, left_on = 'customer_id', right_on = 'customer_id', how = 'inner')
df_CI_S.head()
```

```
[111]: customer_id  age    long_lat  gender  INTER BANK  PAY/SALARY  PAYMENT  \
0  CUS-1005756958  53  153.03 -27.51    F           0  12616.11  1296
1  CUS-1117979751  21  115.81 -31.82    M        1001  25050.55  3957
2  CUS-1140341822  28  144.97 -37.42    M        270  11499.06  852
3  CUS-1147642491  34  151.04 -33.77    F        250  22248.07  2597
```

4	CUS-1196156254	34	138.52	-35.01	F	767	27326.11	3017
---	----------------	----	--------	--------	---	-----	----------	------

	PHONE BANK	POS	SALES-POS	Salary_Count
0	546	748.01	1062.85	13
1	0	606.37	3369.45	7
2	0	1356.47	3033.07	6
3	0	2062.31	1823.44	13
4	0	2305.65	2634.96	7

```
[112]: # Calculate annual salary
# if Salary_Count >= 12 then payment is weekly
# if Salary_Count <= 5 then payment is monthly
# if other value then payment is fortnightly
```

```
[113]: df_CI_S ['Annual_Salary'] = 0
for i in range (0, len(df_CI_S)):
    if int(df_CI_S.Salary_Count[i]) >= 12:
        df_CI_S.Annual_Salary[i] = df_CI_S ['PAY/SALARY'][i]/(df_CI_S.
→Salary_Count[i]) / 7 * 356
    elif int(df_CI_S.Salary_Count[i]) <= 5:
        df_CI_S.Annual_Salary[i] = df_CI_S ['PAY/SALARY'][i]/(df_CI_S.
→Salary_Count[i]) * 12
    else:
        df_CI_S.Annual_Salary[i] = df_CI_S ['PAY/SALARY'][i]/(df_CI_S.
→Salary_Count[i]) / 14 * 356

# all transaction multiply 4 to obtain for one year

df_CI_S [['INTER BANK', 'PAYMENT', 'PHONE BANK', 'POS',
→'SALES-POS']] =4 * df_CI_S [['INTER BANK', 'PAYMENT', 'PHONE BANK', 'POS',
→'SALES-POS']]

df_CI_S.head()
```

C:\Users\atefeh\Anaconda3\lib\site-packages\ipykernel\_launcher.py:4:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

after removing the cwd from sys.path.

C:\Users\atefeh\Anaconda3\lib\site-packages\ipykernel\_launcher.py:8:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

C:\Users\atefeh\Anaconda3\lib\site-packages\ipykernel\_launcher.py:6:

SettingWithCopyWarning:

A value is trying to be set on a copy of a slice from a DataFrame

See the caveats in the documentation: <http://pandas.pydata.org/pandas-docs/stable/indexing.html#indexing-view-versus-copy>

```
[113]:      customer_id  age      long_lat  gender  INTER BANK  PAY/SALARY  PAYMENT  \
0  CUS-1005756958   53   153.03 -27.51      F         0    12616.11     5184
1  CUS-1117979751   21   115.81 -31.82      M    4004    25050.55    15828
2  CUS-1140341822   28   144.97 -37.42      M    1080    11499.06     3408
3  CUS-1147642491   34   151.04 -33.77      F    1000    22248.07    10388
4  CUS-1196156254   34   138.52 -35.01      F    3068    27326.11    12068
```

```
      PHONE BANK      POS  SALES-POS  Salary_Count  Annual_Salary
0      2184  2992.04    4251.40         13      49355
1         0  2425.48    13477.80          7      90999
2         0  5425.88    12132.28          6      48734
3         0  8249.24     7293.76         13      87036
4         0  9222.60    10539.84          7      99266
```

```
[114]: # find the initial balance of each customer
df_B = dataset[['customer_id', 'balance']]
df_CI_B = df_B.drop_duplicates(subset = ['customer_id'], keep = 'first')

# merge with previous dataframe
df_S_B = df_CI_S.merge(df_CI_B , left_on = 'customer_id', right_on = 'customer_id', how = 'inner')

# adding a column for total spend
df_S_B['total spend'] = df_S_B['INTER BANK'] + df_S_B['PAYMENT'] + df_S_B['PHONE BANK'] + df_S_B['POS'] + df_S_B['SALES-POS']

df_S_B.head()
```

```
[114]:      customer_id  age      long_lat  gender  INTER BANK  PAY/SALARY  PAYMENT  \
0  CUS-1005756958   53   153.03 -27.51      F         0    12616.11     5184
1  CUS-1117979751   21   115.81 -31.82      M    4004    25050.55    15828
2  CUS-1140341822   28   144.97 -37.42      M    1080    11499.06     3408
3  CUS-1147642491   34   151.04 -33.77      F    1000    22248.07    10388
4  CUS-1196156254   34   138.52 -35.01      F    3068    27326.11    12068
```

```
      PHONE BANK      POS  SALES-POS  Salary_Count  Annual_Salary  balance  \
0      2184  2992.04    4251.40         13      49355     463.96
1         0  2425.48    13477.80          7      90999     2335.35
2         0  5425.88    12132.28          6      48734      823.53
```

3	0	8249.24	7293.76	13	87036	1726.28
4	0	9222.60	10539.84	7	99266	12529.59

	total spend
0	14611.44
1	35735.28
2	22046.16
3	26931.00
4	34898.44

```
[115]: # find state of each customer base on longitude and latitude information
"""
df_S_B[['long', 'lat']] = df_S_B.long_lat.str.split(' ', expand = True)
df_S_B['lat_long'] = df_S_B['lat'].map(str) + ',' + df_S_B['long'].map(str)

import geopy
from geopy.geocoders import Nominatim
def LoctoState(LOC):
    locator = Nominatim(user_agent="myGeocoder")
    coordinates = LOC
    location = locator.reverse(coordinates, exactly_one = True, timeout = 10)
    address = location.raw['address']
    state = address.get('state')
    return state
#print (LoctoState( "-27.51,153.03"))

df_S_B['state'] = ''
for i in range (0, len(df_S_B)):
    df_S_B['state'][i] = LoctoState(df_S_B['lat_long'][i])

df_S_B.head()
"""
```

```
[115]: '\ndf_S_B[['long', 'lat']] = df_S_B.long_lat.str.split(' ', expand =
True)\ndf_S_B['lat_long'] = df_S_B['lat'].map(str) + ',' +
df_S_B['long'].map(str)\n\nimport geopy\nfrom geopy.geocoders import
Nominatim\ndef LoctoState(LOC):\n    locator =
Nominatim(user_agent="myGeocoder")\n    coordinates = LOC\n    location =
locator.reverse(coordinates, exactly_one = True, timeout = 10)\n    address =
location.raw['address']\n    state = address.get('state')\n    return
state\n#print (LoctoState( "-27.51,153.03"))\n\ndf_S_B['state'] = ''\nfor i
in range (0, len(df_S_B)):\n    df_S_B['state'][i] =
LoctoState(df_S_B['lat_long'][i])\n    \ndf_S_B.head()\n'
```

```
[116]: df_F = df_S_B.drop(['customer_id', 'long_lat', 'PAY/SALARY', 'Salary_Count'],
→axis = 1)
df_F.head()
```


```
[116]:
```

	age	gender	INTER BANK	PAYMENT	PHONE BANK	POS	SALES-POS	\
0	53	F	0	5184	2184	2992.04	4251.40	
1	21	M	4004	15828	0	2425.48	13477.80	
2	28	M	1080	3408	0	5425.88	12132.28	
3	34	F	1000	10388	0	8249.24	7293.76	
4	34	F	3068	12068	0	9222.60	10539.84	

	Annual_Salary	balance	total spend
0	49355	463.96	14611.44
1	90999	2335.35	35735.28
2	48734	823.53	22046.16
3	87036	1726.28	26931.00
4	99266	12529.59	34898.44

```
[117]: # Histogram of Data visualisation
df_F.hist(rwidth = 0.9)
plt.tight_layout()
```




```
[118]: # Data visualisation
# Visualise the continuous features vs PAY/SALARY
#plt.subplot(8,1,1)
plt.title('INTER BANK vs Annual_Salary')
plt.scatter(df_F['INTER BANK'], df_F['Annual_Salary'], s = 2)
plt.xlabel('INTER BANK')
```

```
plt.ylabel('Annual_Salary')
```

```
#plt.tight_layout()
```


[118]: Text(0, 0.5, 'Annual\_Salary')



```
[119]: #plt.subplot(8,1,2)
plt.title('PAYMENT vs Annual_Salary')
plt.scatter(df_F['PAYMENT'], df_F['Annual_Salary'], s = 2)
plt.xlabel('PAYMENT')
plt.ylabel('Annual_Salary')
```


[119]: Text(0, 0.5, 'Annual\_Salary')






```
[120]: #plt.subplot(8,1,3)
plt.title('PHONE BANK vs Annual_Salary')
plt.scatter(df_F['PHONE BANK'], df_F['Annual_Salary'], s = 2)
plt.xlabel('PHONE BANK')
plt.ylabel('Annual_Salary')
```

```
[120]: Text(0, 0.5, 'Annual_Salary')
```




```
[121]: #plt.subplot(8,1,4)
plt.title('POS vs Annual_Salary')
plt.scatter(df_F['POS'], df_F['Annual_Salary'], s = 2)
plt.xlabel('POS')
plt.ylabel('Annual_Salary')
```

```
[121]: Text(0, 0.5, 'Annual_Salary')
```




```
[122]: #plt.subplot(8,1,5)
plt.title('SALES-POS vs Annual_Salary')
plt.scatter(df_F['SALES-POS'], df_F['Annual_Salary'], s = 2)
plt.xlabel('SALES-POS')
plt.ylabel('Annual_Salary')
```

```
[122]: Text(0, 0.5, 'Annual_Salary')
```




```
[123]: #plt.subplot(8,1,6)
plt.title('age vs Annual_Salary')
plt.scatter(df_F['age'], df_F['Annual_Salary'], s = 2)
plt.xlabel('age')
plt.ylabel('Annual_Salary')
```

```
[123]: Text(0, 0.5, 'Annual_Salary')
```



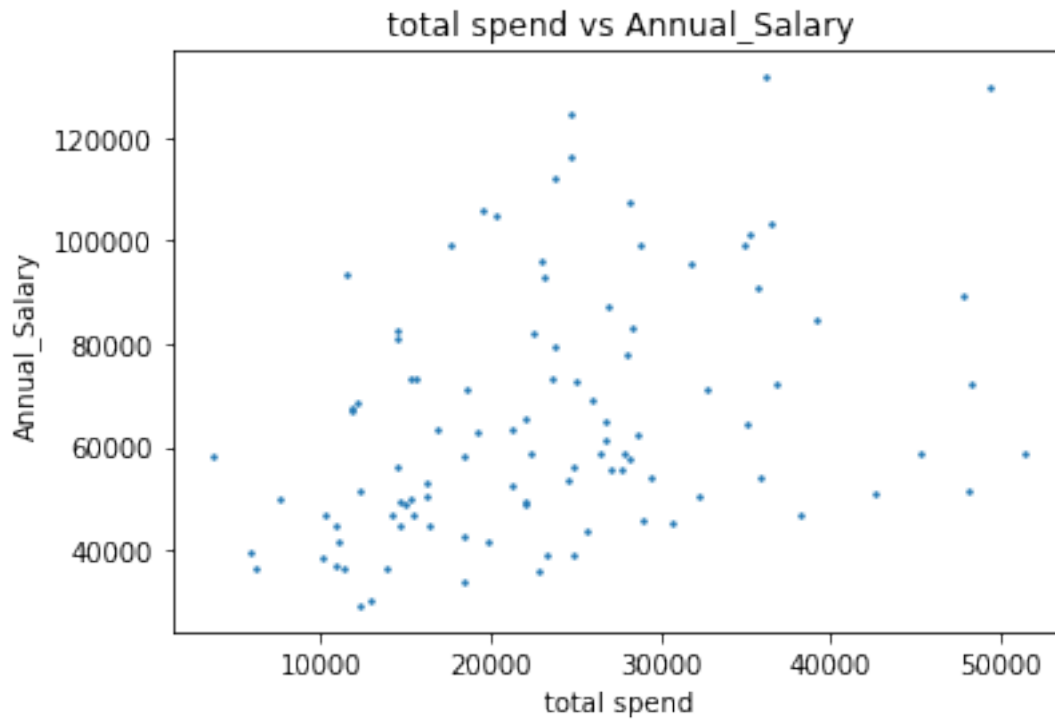
```
[124]: #plt.subplot(8,1,7)
plt.title('Balance vs Annual_Salary')
plt.scatter(df_F['balance'], df_F['Annual_Salary'], s = 2)
plt.xlabel('balance')
plt.ylabel('Annual_Salary')
```

```
[124]: Text(0, 0.5, 'Annual_Salary')
```



```
[125]: #plt.subplot(8,1,8)
plt.title('total spend vs Annual_Salary')
plt.scatter(df_F['total spend'], df_F['Annual_Salary'], s = 2)
plt.xlabel('total spend')
plt.ylabel('Annual_Salary')
```


[125]: Text(0, 0.5, 'Annual\_Salary')



```
[126]: # Data visulisation
# Visualise the categorical feature vs PAY/SALARY

cat_list = df_F['gender'].unique()
cat_average = df_F.groupby('gender')['Annual_Salary'].mean()
#cat_average = df_F.groupby('gender').mean()['Annual_Salary']
plt.title('Average Annual_Salary vs gender')
plt.bar(cat_list, cat_average)
```

[126]: <BarContainer object of 2 artists>



```
[127]: # Check for outlier
df_F['Annual_Salary'].describe()
#df_F['Annual_Salary'].quantile([0.05, 0.1, 0.15, 0.9, 0.95, 0.99])
```

```
[127]: count      100.000000
mean      65352.270000
std       23633.324404
min       29293.000000
25%       48267.500000
50%       58803.000000
75%       79821.250000
max       131618.000000
Name: Annual_Salary, dtype: float64
```

```
[128]: # Check linearity using correlation coefficient matrix
correlation = df_F[['Annual_Salary', 'age', 'INTER BANK', 'PAYMENT', 'PHONE BANK',
                    'POS', 'SALES-POS',
                    'balance', 'total spend']].corr()
print(correlation)
```

	Annual_Salary	age	INTER BANK	PAYMENT	PHONE BANK	\
Annual_Salary	1.000000	-0.061377	0.352362	0.639631	0.033414	
age	-0.061377	1.000000	-0.099233	0.026884	0.103961	
INTER BANK	0.352362	-0.099233	1.000000	0.087386	-0.081680	
PAYMENT	0.639631	0.026884	0.087386	1.000000	-0.132095	
PHONE BANK	0.033414	0.103961	-0.081680	-0.132095	1.000000	



POS	-0.086938	-0.036929	0.181437	-0.123618	-0.052313
SALES-POS	0.100400	-0.139284	0.158792	0.121610	0.007496
balance	0.110321	0.237992	0.211241	0.018268	0.026537
total spend	0.371378	-0.086176	0.476295	0.416710	0.015633

	POS	SALES-POS	balance	total spend
Annual_Salary	-0.086938	0.100400	0.110321	0.371378
age	-0.036929	-0.139284	0.237992	-0.086176
INTER BANK	0.181437	0.158792	0.211241	0.476295
PAYMENT	-0.123618	0.121610	0.018268	0.416710
PHONE BANK	-0.052313	0.007496	0.026537	0.015633
POS	1.000000	0.418105	-0.000239	0.689232
SALES-POS	0.418105	1.000000	-0.153300	0.756790
balance	-0.000239	-0.153300	1.000000	0.002437
total spend	0.689232	0.756790	0.002437	1.000000

```
[129]: # Drop irrelevant features
df_ss = df_F.drop(['age', 'PHONE BANK', 'POS'], axis = 1)
df_ss.head()
```

```
[129]:  gender  INTER BANK  PAYMENT  SALES-POS  Annual_Salary  balance  total spend
0      F           0      5184      4251.40           49355      463.96      14611.44
1      M          4004      15828      13477.80           90999      2335.35      35735.28
2      M          1080       3408      12132.28           48734       823.53      22046.16
3      F          1000      10388       7293.76           87036      1726.28      26931.00
4      F          3068      12068      10539.84           99266     12529.59      34898.44
```

```
[130]: # Create dummy variable
df_ss['gender'].dtype
df_sm = pd.get_dummies(df_ss, drop_first = True)
#df_sm = df_ss
```

```
[131]: # Normalise the data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df_sm['INTER BANK'] = scaler.fit_transform(df_sm[['INTER BANK']])
df_sm['PAYMENT'] = scaler.fit_transform(df_sm[['PAYMENT']])
df_sm['SALES-POS'] = scaler.fit_transform(df_sm[['SALES-POS']])
df_sm['balance'] = scaler.fit_transform(df_sm[['balance']])
df_sm['total spend'] = scaler.fit_transform(df_sm[['total spend']])
df_sm['Annual_Salary'] = scaler.fit_transform(df_sm[['Annual_Salary']])
```

```
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:645:
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
```

```
    return self.partial_fit(X, y)
```

```
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\base.py:464:
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
```

```

    return self.fit(X, **fit_params).transform(X)
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:645:
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
    return self.partial_fit(X, y)
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\base.py:464:
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
    return self.fit(X, **fit_params).transform(X)
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:645:
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
    return self.partial_fit(X, y)
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\base.py:464:
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
    return self.fit(X, **fit_params).transform(X)

```

```

[132]: # Split into train and test
from sklearn.model_selection import train_test_split

X = df_sm.drop(['Annual_Salary'], axis = 1)
Y = df_sm[['Annual_Salary']]

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
→random_state = 1234)

```

## 1 Linear Regression

```

[133]: # Fit and score the model
# Linear regression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error

LR = LinearRegression()
LR.fit(X_train, Y_train)
Y_Predicted = LR.predict(X_test)

score_train = LR.score(X_train, Y_train)
score_test = LR.score(X_test, Y_test)
#r2_score = r2_score(Y_test, Y_Predicted)
RMSE = math.sqrt(mean_squared_error(Y_test, Y_Predicted))

print('r2 score for Linear Regression is:' ,score_test)

```

```
print('RMSE for Linear Regression is:' ,RMSE)
```

r2 score for Linear Regression is: 0.5763367529452759

RMSE for Linear Regression is: 0.6476652932203193

```
[134]: # cross_val_score
from sklearn.model_selection import cross_val_score
scores = cross_val_score(LR, X_train, Y_train, scoring='r2', cv=5)
scores
```

```
[134]: array([ 0.2288959 , -0.13826055,  0.34843882,  0.44948336, -0.89811154])
```

```
[135]: # KFold
from sklearn.model_selection import KFold
folds = KFold(n_splits = 5, shuffle = True, random_state = 100)
scores = cross_val_score(LR, X_train, Y_train, scoring='r2', cv=folds)
scores
```

```
[135]: array([ 0.2989966 ,  0.64696935,  0.57678064, -0.17722345, -0.11800458])
```

## 2 Feature Selection

```
[136]: # feature selection
from sklearn.feature_selection import f_regression
result = f_regression(X,Y)
f_score = result [0]
p_value = result [1]

columns = list(X.columns)
print("  Feature  ", "          F_score  ", "          P_value  ")
print("  -----  ", "          -----  ", "          -----  ")
for i in range(0, len(columns)):
    f1= f_score[i]
    p1= p_value[i]
    print(" ", columns[i].ljust(10), " ", f1, " ", p1)
```

Feature	F_score	P_value
-----	-----	-----
INTER BANK	13.892456810874593	0.00032395470010236047
PAYMENT	67.85635791948742	7.878230949637538e-13
SALES-POS	0.997918573516517	0.32027489509276974
balance	1.2074184722859997	0.2745362021940989
total spend	15.678767795365692	0.00014233394008753423
gender_M	1.9135367414106539	0.16971417787004986

C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\utils\validation.py:761:  
DataConversionWarning: A column-vector y was passed when a 1d array was

expected. Please change the shape of y to (n\_samples, ), for example using ravel().

```
y = column_or_1d(y, warn=True)
```

### 3 Linear Regression after FS

```
[137]: # Drop irrelevant features
df_ss = df_F.drop([ 'age', 'PHONE BANK', 'POS', 'balance', 'SALES-POS'], axis = 1)
df_ss.head()
```

```
[137]:
```

	gender	INTER BANK	PAYMENT	Annual_Salary	total spend
0	F	0	5184	49355	14611.44
1	M	4004	15828	90999	35735.28
2	M	1080	3408	48734	22046.16
3	F	1000	10388	87036	26931.00
4	F	3068	12068	99266	34898.44

```
[138]: # Create dummy variable
df_ss['gender'].dtype
df_sm = pd.get_dummies(df_ss, drop_first = True)
```

```
[139]: # Normalise the data
from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
df_sm['INTER BANK'] = scaler.fit_transform(df_sm[['INTER BANK']])
df_sm['PAYMENT'] = scaler.fit_transform(df_sm[['PAYMENT']])
df_sm['total spend'] = scaler.fit_transform(df_sm[['total spend']])
df_sm['Annual_Salary'] = scaler.fit_transform(df_sm[['Annual_Salary']])
```

```
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:645:
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
```

```
return self.partial_fit(X, y)
```

```
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\base.py:464:
```

```
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
```

```
return self.fit(X, **fit_params).transform(X)
```

```
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:645:
```

```
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
```

```
return self.partial_fit(X, y)
```

```
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\base.py:464:
```

```
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
```

```
return self.fit(X, **fit_params).transform(X)
```

```
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\preprocessing\data.py:645:
```

```
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
```

```

    return self.partial_fit(X, y)
C:\Users\atefeh\Anaconda3\lib\site-packages\sklearn\base.py:464:
DataConversionWarning: Data with input dtype int64 were all converted to float64
by StandardScaler.
    return self.fit(X, **fit_params).transform(X)

```

```

[140]: # Split into train and test
from sklearn.model_selection import train_test_split

X = df_sm.drop(['Annual_Salary'], axis = 1)
Y = df_sm[['Annual_Salary']]

X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size = 0.2,
→random_state = 1234)

```

```

[141]: # Fit and score the model
# Linear regression
from sklearn.linear_model import LinearRegression
from sklearn.metrics import r2_score
from sklearn.metrics import mean_squared_error

LR = LinearRegression()
LR.fit(X_train, Y_train)
Y_Predicted = LR.predict(X_test)

score_train = LR.score(X_train, Y_train)
score_test = LR.score(X_test, Y_test)
RMSE = math.sqrt(mean_squared_error(Y_test, Y_Predicted))

print('r2 score for Linear Regression is:', score_test)
print('RMSE for Linear Regression is:', RMSE)

```

```

r2 score for Linear Regression is: 0.6173311697006681
RMSE for Linear Regression is: 0.6155336134400503

```

```

[142]: # cross_val_score
from sklearn.model_selection import cross_val_score
scores = cross_val_score(LR, X_train, Y_train, scoring='r2', cv=5)
scores

```

```

[142]: array([ 0.23932364, -0.12905291,  0.44972776,  0.49994856, -0.49757349])

```

```

[143]: # KFold
from sklearn.model_selection import KFold
folds = KFold(n_splits = 5, shuffle = True, random_state = 100)
scores = cross_val_score(LR, X_train, Y_train, scoring='r2', cv=folds)
scores

```

```

[143]: array([ 0.40019539,  0.68358956,  0.58267536, -0.19175429, -0.09101695])

```

[ ]:	
[ ]:	